

customer-segmentation

October 18, 2023

1 Customer Segmentation

1.1 Table of Contents

1. **Functions**
2. Preprocess Data
 - 2.1 Load Data
 - 2.2 NaN Values (i.e., " ?")
 - 2.3 Preprocessing Steps
3. **EDA**
4. **Clustering**
 - 4.1 k-Means
 - 4.2 Normalized k-Means
 - 4.3 **DBSCAN**
5. **Visualization**
 - 5.1 **PCA**
 - 5.2 t-SNE
 - 5.3 3D Animation
6. **Evaluation**
7. What makes a cluster unique?
 - 7.1 **Variance**
 - 7.2 Feature Importance

1.2 1. Functions

[Back to Table of Contents](#)

```

[2]: # Data handling
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

# Clustering
from sklearn.cluster import KMeans, DBSCAN
from sklearn import preprocessing
from sklearn.metrics import silhouette_score

# Dimensionality reduction
from sklearn.manifold import TSNE
from sklearn.decomposition import PCA

# Visualization
import seaborn as sns
import plotly.express as px
import matplotlib.pyplot as plt
import mpl_toolkits.mplot3d.axes3d as p3
from matplotlib import animation

%matplotlib inline

def load_preprocess_data():
    """ Load and preprocess data
    """

    # Load data
    df = pd.read_csv("data.csv")

    # remove empty values
    df = df.loc[df.TotalCharges!=" ", :]
    df.TotalCharges = df.TotalCharges.astype(float)

    # Label data correctly
    replace_cols = [ 'OnlineSecurity', 'OnlineBackup', 'DeviceProtection',
                     'TechSupport', 'StreamingTV', 'StreamingMovies', 'Partner',
    ↪ 'Dependents',
                     'PhoneService', 'MultipleLines', 'PaperlessBilling', 'Churn']
    for i in replace_cols :
        df.loc[:, i] = df.loc[:, i].replace({'No internet service' : 'No', 'No
    ↪ phone service': 'No'})
        df.loc[:, i] = df.loc[:, i].map({'No':0, 'Yes':1})
    df.gender = df.gender.map({"Female":0, "Male":1})

    # One-hot encoding of variables
    others_categorical = ['Contract', 'PaymentMethod', 'InternetService']

```

```

for i in others_categorical:
    df = df.join(pd.get_dummies(df[i], prefix=i))
df.drop(others_categorical, axis=1, inplace=True)

# Calculate number of services
services = ['PhoneService', 'MultipleLines', 'OnlineSecurity',
            'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
            'StreamingMovies', 'InternetService_DSL', 'InternetService_Fiber_
↳optic',
            'InternetService_No']
df['nr_services'] = df.apply(lambda row: sum([row[x] for x in services[:
↳-1]]), 1)

return df.drop('customerID', 1)

def plot_corr(df):
    corr = df.corr()
    mask = np.zeros_like(corr, dtype=np.bool)
    mask[np.triu_indices_from(mask)] = True
    f, ax = plt.subplots(figsize=(11, 9))
    cmap = sns.diverging_palette(220, 10, as_cmap=True)
    sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,
                square=True, linewidths=.5, cbar_kws={"shrink": .5})

def plot_tsne(tsne_data, kmeans_labels):
    df_tsne = pd.DataFrame(tsne_data).rename({0: 'x', 1: 'y'}, axis=1)
    df_tsne['z'] = kmeans_labels
    sns.scatterplot(x=df_tsne.x, y=df_tsne.y, hue=df_tsne.z, palette="Set2")
    plt.show()

def prepare_pca(n_components, data, kmeans_labels):
    names = ['x', 'y', 'z']
    matrix = PCA(n_components=n_components).fit_transform(data)
    df_matrix = pd.DataFrame(matrix)
    df_matrix.rename({i:names[i] for i in range(n_components)}, axis=1,
↳inplace=True)
    df_matrix['labels'] = kmeans_labels

    return df_matrix

def prepare_tsne(n_components, data, kmeans_labels):
    names = ['x', 'y', 'z']
    matrix = TSNE(n_components=n_components).fit_transform(data)
    df_matrix = pd.DataFrame(matrix)
    df_matrix.rename({i:names[i] for i in range(n_components)}, axis=1,
↳inplace=True)
    df_matrix['labels'] = kmeans_labels

```

```

    return df_matrix

def plot_3d(df, name='labels'):
    iris = px.data.iris()
    fig = px.scatter_3d(df, x='x', y='y', z='z',
                        color=name, opacity=0.5)

    fig.update_traces(marker=dict(size=3))
    fig.show()

def plot_animation(df, label_column, name):
    def update(num):
        ax.view_init(200, num)

    N=360
    fig = plt.figure()
    ax = fig.add_subplot(111, projection='3d')
    ax.scatter(tsne_3d_df['x'], tsne_3d_df['y'], tsne_3d_df['z'],
               c=tsne_3d_df[label_column],
               s=6, depthshade=True, cmap='Paired')
    ax.set_zlim(-15, 25)
    ax.set_xlim(-20, 20)
    plt.tight_layout()
    ani = animation.FuncAnimation(fig, update, N, blit=False, interval=50)
    ani.save('{}.gif'.format(name), writer='imagemagick')
    plt.show()

```

1.3 2. Preprocess Data

[Back to Table of Contents](#)

Demographic * Gender * SeniorCitizen * Partner * Dependents * Tenure

Services * PhoneService * MultipleLines * InternetService * OnlineSecurity * OnlineBackup *
DeviceProtection * TechSupport * StreamingTV * StreamingMovies

Customer account information * Contract * PaperlessBilling * PaymentMethod * MonthlyCharges
* TotalCharges

Target * Churn

No = 0 Yes = 1

Female = 0 Male = 1


```
[3]: df = load_preprocess_data()
```

1.4 3. EDA


[Back to Table of Contents](#)

```
[10]: sns.scatterplot(df.TotalCharges, df.tenure, df.nr_services)
```

```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1e91517c4e0>
```



```
[11]: plot_corr(df)
```



1.5 4. Clustering

[Back to Table of Contents](#)


```
[4]: df = df.drop(["Churn"], 1)
```

1.5.1 4.1. k-Means

[Back to Table of Contents](#)

```
[13]: scores = [KMeans(n_clusters=i+2).fit(df).inertia_ for i in range(10)]
sns.lineplot(np.arange(2, 12), scores)
plt.xlabel('Number of clusters')
plt.ylabel("Inertia")
plt.title("Inertia of k-Means versus number of clusters")
```

```
[13]: Text(0.5, 1.0, 'Inertia of k-Means versus number of clusters')
```



```
[14]: kmeans = KMeans(n_clusters=4)
      kmeans.fit(df)
```


```
[14]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=4, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=None, tol=0.0001, verbose=0)
```

1.5.2 4.2. Normalized k-Means

[Back to Table of Contents](#)

```
[15]: normalized_vectors = preprocessing.normalize(df)
      scores = [KMeans(n_clusters=i+2).fit(normalized_vectors).inertia_ for i in
               ↪range(10)]
      sns.lineplot(np.arange(2, 12), scores)
      plt.xlabel('Number of clusters')
      plt.ylabel("Inertia")
      plt.title("Inertia of Cosine k-Means versus number of clusters")
```

```
[15]: Text(0.5, 1.0, 'Inertia of Cosine k-Means versus number of clusters')
```



```
[16]: normalized_kmeans = KMeans(n_clusters=4)
      normalized_kmeans.fit(normalized_vectors)
```

```
[16]: KMeans(algorithm='auto', copy_x=True, init='k-means++', max_iter=300,
            n_clusters=4, n_init=10, n_jobs=None, precompute_distances='auto',
            random_state=None, tol=0.0001, verbose=0)
```

1.5.3 4.3. DBSCAN

[Back to Table of Contents](#)

```
[18]: min_samples = df.shape[1]+1 # Rule of thumb; number of dimensions D in the
      ↪ data set, as minPts D + 1
      dbscan = DBSCAN(eps=3.5, min_samples=min_samples).fit(df)
```

1.6 5. Visualization


[Back to Table of Contents](#)

1.6.1 5.1. PCA

[Back to Table of Contents](#)

```
[112]: pca_df = prepare_pca(3, df, normalized_kmeans.labels_)
      sns.scatterplot(x=pca_df.x, y=pca_df.y, hue=pca_df.labels, palette="Set2")
```


[112]: <matplotlib.axes._subplots.AxesSubplot at 0x2292d211e48>



```
[108]: pca_df = prepare_pca(3, df, normalized_kmeans.labels_)
plot_3d(pca_df)
```

1.6.2 5.2. t-SNE

[Back to Table of Contents](#)

```
[21]: tsne_3d_df = prepare_tsne(3, df, kmeans.labels_)
```

```
[22]: plot_3d(tsne_3d_df)
```

```
[307]: tsne_3d_df['normalized_kmeans'] = normalized_kmeans.labels_
plot_3d(tsne_3d_df, name='normalized_kmeans')
```


```
[67]: min_samples = clustering.shape[1]+1 # Rule of thumb; number of dimensions D in
    ↪ the data set, as minPts D + 1
dbscan = DBSCAN(eps=3.5, min_samples=min_samples).fit(clustering)
# dbscan = DBSCAN(eps=50, min_samples=min_samples).fit(clustering)
tsne_3d_df['dbscan'] = [str(label) for label in dbscan.labels_]
plot_3d(tsne_3d_df, name='dbscan')
```

1.6.3 5.3. 3D Animation

[Back to Table of Contents](#)

```
[321]: tsne_3d_df.dbscan = tsne_3d_df.dbscan.astype(int)
       plot_animation(tsne_3d_df, 'normalized_kmeans', 'normalized_kmeans_new')
```

MovieWriter imagemagick unavailable; trying to use <class 'matplotlib.animation.PillowWriter'> instead.



1.7 6. Evaluation

[Back to Table of Contents](#)

```
[7]: kmeans = KMeans(n_clusters=4).fit(df)

normalized_vectors = preprocessing.normalize(df)
normalized_kmeans = KMeans(n_clusters=4).fit(normalized_vectors)

min_samples = df.shape[1]+1 # Rule of thumb; number of dimensions D in the
    ↳ data set, as minPts D + 1
dbscan = DBSCAN(eps=3.5, min_samples=min_samples).fit(df)
```

```
[105]: print('kmeans: {}'.format(silhouette_score(df, kmeans.labels_,
    ↳ metric='euclidean')))
```

```
print('Cosine kmeans: {}'.format(silhouette_score(normalized_vectors,
↪normalized_kmeans.labels_, metric='cosine')))
print('DBSCAN: {}'.format(silhouette_score(df, dbscan.labels_,
↪metric='cosine')))
```

kmeans: 0.6018318118002677
Cosine kmeans: 0.8633823077551214
DBSCAN: 0.8302013261718773


1.8 7. What makes a cluster unique?

[Back to Table of Contents](#)

One way to see the differences between clusters is to take the average value of each cluster and visualize it.

```
[9]: # Setting all variables between 0 and 1 in order to better visualize the results
scaler = MinMaxScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df))
df_scaled.columns = df.columns
df_scaled['dbscan'] = dbscan.labels_
```

```
[17]: # df = load_preprocess_data()
df['dbscan'] = dbscan.labels_
tidy = df_scaled.melt(id_vars='dbscan')
fig, ax = plt.subplots(figsize=(15, 5))
sns.barplot(x='dbscan', y='value', hue='variable', data=tidy, palette='Set3')
plt.legend([''])
# plt.savefig("mess.jpg", dpi=300)
plt.savefig("dbscan_mess.jpg", dpi=300)
```



The problem with this approach is that we simply have too many variables. Not all of them are likely to be important when creating the clusters. Instead, I will select the most important columns based on the following approach:

1.8.1 7.1. Variance

[Back to Table of Contents](#)

What I essentially do is group datapoints by cluster and take the average. Then, I calculate the standard deviation between those values for each variable. Variables with a high standard deviation indicate that there are large differences between clusters and that the variable might be important.

```
[11]: df_mean = df_scaled.loc[df_scaled.dbscan!=-1, :].groupby('dbscan').mean().  
      ↪reset_index()
```

```
[12]: df_mean
```

```
[12]:
```

	dbscan	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	0	0.593750	0.018750	0.168750	0.218750	0.0	
1	1	0.609756	0.105691	0.097561	0.154472	0.0	
2	2	0.480447	0.251397	0.195531	0.145251	0.0	
3	3	0.432432	0.216216	0.108108	0.189189	0.0	

	PhoneService	MultipleLines	OnlineSecurity	OnlineBackup	...	\
0	1.000000	0.000000	0.000000	0.000000	...	
1	0.967480	0.073171	0.130081	0.089431	...	
2	1.000000	0.184358	0.022346	0.083799	...	
3	0.162162	0.162162	0.000000	0.000000	...	

	Contract_One year	Contract_Two year	\
0	0.031250	0.0125	
1	0.000000	0.0000	
2	0.000000	0.0000	
3	0.027027	0.0000	

	PaymentMethod_Bank transfer (automatic)	\
0	0.081250	
1	0.081301	
2	0.039106	
3	0.054054	

	PaymentMethod_Credit card (automatic)	PaymentMethod_Electronic check	\
0	0.068750	0.112500	
1	0.073171	0.398374	
2	0.027933	0.731844	
3	0.000000	0.432432	

	PaymentMethod_Mailed check	InternetService_DSL	\
0	0.737500	0.000000	
1	0.447154	1.000000	
2	0.201117	0.000000	
3	0.513514	0.837838	


	InternetService_Fiber optic	InternetService_No	nr_services
0	0.0	1.000000	0.000000
1	0.0	0.000000	0.183943
2	1.0	0.000000	0.194134
3	0.0	0.162162	0.020270

[4 rows x 28 columns]

```
[15]: # Setting all variables between 0 and 1 in order to better visualize the results
# df = load_preprocess_data().drop("Churn", 1)
scaler = MinMaxScaler()
df_scaled = pd.DataFrame(scaler.fit_transform(df))
df_scaled.columns = df.columns
df_scaled['dbscan'] = dbscan.labels_

# Calculate variables with largest differences (by standard deviation)
# The higher the standard deviation in a variable based on average values for
↳ each cluster
# The more likely that the variable is important when creating the cluster
df_mean = df_scaled.loc[df_scaled.dbscan!=-1, :].groupby('dbscan').mean().
↳ reset_index()
results = pd.DataFrame(columns=['Variable', 'Std'])
for column in df_mean.columns[1:]:
    results.loc[len(results), :] = [column, np.std(df_mean[column])]
selected_columns = list(results.sort_values('Std', ascending=False).head(7).
↳ Variable.values) + ['dbscan']

# Plot data
tidy = df_scaled[selected_columns].melt(id_vars='dbscan')
fig, ax = plt.subplots(figsize=(15, 5))
sns.barplot(x='dbscan', y='value', hue='variable', data=tidy, palette='Set3')
plt.legend(loc='upper right')
plt.savefig("dbscan_results.jpg", dpi=300)
```




1.8.2 7.2. Feature Importance

[Back to Table of Contents](#)

```
[18]: from sklearn.ensemble import RandomForestClassifier
y = df.iloc[:, -1]
X = df.iloc[:, :-1]
clf = RandomForestClassifier(n_estimators=100).fit(X, y)
selected_columns = list(pd.DataFrame(np.array([clf.feature_importances_, X.
↳ columns]).T, columns=['Importance', 'Feature'])
    .sort_values("Importance", ascending=False)
    .head(7)
    .Feature
    .values)

[20]: # Plot data
tidy = df_scaled[selected_columns+['dbscan']].melt(id_vars='dbscan')
fig, ax = plt.subplots(figsize=(15, 5))
sns.barplot(x='dbscan', y='value', hue='variable', data=tidy, palette='Set3')
plt.legend(loc='upper right')
plt.savefig('randomforest.jpg', dpi=300)
```



```
[ ]:
```